Advanced Data Analysis Techniques

Project: Daily Temperature Forecasting Using Time Series Analysis (ARIMA)

Name: Vidhya P

Internship name: Data Science Internship

Submitted by: 02-07-2026

Acknowledgement

I would like to express my sincere gratitude to my internship mentor and the organization for providing me with the opportunity to work on this project as part of the Data Science Internship. This project has helped me gain practical knowledge of advanced data analysis techniques and strengthened my understanding of time series forecasting using Python.

I would also like to thank my faculty members for their continuous encouragement and valuable guidance throughout the project. Their suggestions and support helped me complete the project successfully.

Finally, I extend my gratitude to everyone who directly or indirectly contributed to the successful completion of this project.

Abstract

Time Series Analysis is one of the most widely used advanced data analysis techniques for analyzing sequential data collected over regular intervals of time. It is extensively applied in domains such as weather forecasting, stock market prediction, sales forecasting, energy demand estimation, healthcare monitoring, and traffic analysis.

This project focuses on forecasting daily minimum temperatures using historical weather data through the ARIMA (AutoRegressive Integrated Moving Average) model. The dataset consists of daily minimum temperature observations collected over a ten-year period. The project begins with data preprocessing, including date conversion, missing value analysis, and statistical exploration. Various visualization techniques are used to identify trends and patterns in the data. Before building the forecasting model, the stationarity of the dataset is verified using the Augmented Dickey-Fuller (ADF) test.

An ARIMA model is then trained using the historical temperature data to predict future temperature values. The forecasting performance is evaluated using standard metrics such as Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE).

Graphical comparisons between actual and predicted temperatures are also presented to assess the model's effectiveness.

The implementation demonstrates the complete workflow of a time series forecasting project using Python libraries such as Pandas, NumPy, Matplotlib, Scikit-learn, and Statsmodels. The project highlights the practical importance of advanced analytical methods in solving real-world prediction problems and provides hands-on experience in developing forecasting models for sequential datasets.

Project Information

Task Name: Advanced Data Analysis Techniques

Chosen Technique: Time Series Analysis

Algorithm Used: ARIMA (AutoRegressive Integrated Moving Average)

Dataset Used: Daily Minimum Temperatures Dataset (1981–1990)

Programming Language: Python

Development Environment: Jupyter Notebook

Libraries Used: Pandas, NumPy, Matplotlib, Statsmodels, Scikit-learn

Project Summary

This project applies the advanced data analysis technique of Time Series Analysis to forecast future daily minimum temperatures using historical weather data. The ARIMA forecasting model is implemented to analyze trends and generate predictions. The project includes data preprocessing, exploratory data analysis, stationarity testing, model building, forecasting, evaluation, and visualization. The implementation demonstrates practical applications of time series forecasting in weather prediction using Python.

Introduction

Time Series Analysis is an advanced data analysis technique used to analyze data collected over regular intervals of time. Unlike traditional datasets where observations are considered independent, time series data contains temporal relationships, meaning that past observations influence future values. This makes time series analysis an essential technique for forecasting and trend analysis.

Time series forecasting has become increasingly important in many industries, including weather prediction, finance, healthcare, agriculture, transportation, energy management, and retail. Organizations use historical data to predict future outcomes, enabling informed decision-making and efficient resource planning.

In this project, historical daily minimum temperature data is analyzed to identify trends and forecast future temperature values. The project utilizes the ARIMA (AutoRegressive Integrated Moving Average) model, a widely used statistical forecasting technique designed specifically for time series data. Before building the forecasting model, the dataset undergoes preprocessing, visualization, and stationarity testing to ensure reliable predictions.

The implementation is carried out using Python and its data science libraries, including Pandas for data manipulation, NumPy for numerical computations, Matplotlib for visualization, Statsmodels for ARIMA modeling, and Scikit-learn for evaluating model performance. The project demonstrates the complete workflow of a real-world time series forecasting application, from data preparation to prediction and evaluation.

Objectives

The primary objectives of this project are:

- To understand the concepts and importance of Time Series Analysis.

- To analyze historical daily minimum temperature data collected over a period of ten years.

- To preprocess and clean the dataset for accurate forecasting.

- To visualize temporal patterns, trends, and variations in temperature data.

- To perform stationarity testing using the Augmented Dickey-Fuller (ADF) test.

- To implement the ARIMA forecasting model for predicting future temperature values.

- To evaluate the forecasting model using performance metrics such as Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE).

- To gain practical experience in applying advanced data analysis techniques using Python.

Problem Statement

Accurate temperature forecasting plays an important role in several real-world applications, including agriculture, disaster management, environmental monitoring, renewable energy planning, and climate research. Traditional forecasting methods often struggle to capture long-term temporal dependencies and changing patterns present in historical weather data.

The objective of this project is to develop a forecasting model capable of predicting future daily minimum temperatures based on historical observations. By applying the ARIMA time series forecasting technique, the project aims to analyze historical trends and generate reliable predictions that can support better planning and decision-making.

Scope of the Project

The scope of this project is limited to forecasting daily minimum temperatures using historical weather data. The project focuses on implementing a univariate time series forecasting model, where only past temperature values are used for prediction.

The project includes data preprocessing, exploratory data analysis, visualization, stationarity testing, model development using ARIMA, forecasting future temperature values, and evaluating model performance. Although the project focuses on weather forecasting, the same methodology can be applied to various domains such as stock market prediction, sales forecasting, energy consumption analysis, traffic prediction, and healthcare monitoring.

This project serves as a practical demonstration of advanced data analysis techniques and provides a strong foundation for solving real-world forecasting problems using time series models.

Dataset Description

Dataset Overview

The dataset used in this project is the **Daily Minimum Temperatures Dataset**, which contains historical weather observations collected in Melbourne, Australia. The dataset records the minimum temperature for each day over a period of ten years, from 1981 to 1990. Since the observations are collected sequentially over time, the dataset is well suited for time series forecasting techniques.

The dataset is publicly available and is widely used for learning and evaluating forecasting algorithms. It provides a realistic example of time-dependent data with seasonal variations and long-term trends.

Dataset Details

Attribute	Description
Dataset Name	Daily Minimum Temperatures
Domain	Weather Forecasting
Number of Records	3650
Number of Features	2
Time Period	1981 – 1990
File Format	CSV

Dataset Attributes

Column Name	Description
Date	Observation date
Temp	Daily minimum temperature in degrees Celsius (°C)

Example Data

Date	Temp
1981-01-01	20.7
1981-01-02	17.9
1981-01-03	18.8
1981-01-04	14.6
1981-01-05	15.8

The "Date" column is converted into datetime format and used as the index of the dataset. The "Temp" column serves as the target variable for forecasting.

Methodology

The project follows a systematic workflow for developing a time series forecasting model. Each stage contributes to building an accurate and reliable prediction model.

Step 1 – Data Collection

The Daily Minimum Temperatures dataset is collected in CSV format and loaded into Python using the Pandas library.

Step 2 – Data Preprocessing

The dataset is cleaned by checking for missing values, converting the Date column into datetime format, and setting it as the index. This prepares the data for time series analysis.

Step 3 – Exploratory Data Analysis (EDA)

Exploratory analysis is performed to understand the characteristics of the data. Line plots, histograms, and box plots are used to visualize temperature distributions and identify trends.

Step 4 – Stationarity Testing

Before applying the ARIMA model, the stationarity of the time series is verified using the Augmented Dickey-Fuller (ADF) test. A stationary time series improves forecasting accuracy.

Step 5 – Model Development

The ARIMA (AutoRegressive Integrated Moving Average) model is trained using historical temperature observations. The model learns the temporal patterns and relationships present in the data.

Step 6 – Forecasting

The trained ARIMA model predicts future daily minimum temperatures based on historical observations.

Step 7 – Model Evaluation

The forecasting model is evaluated using Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE). Lower values indicate better forecasting performance.

Step 8 – Visualization

The actual temperature values are compared with the predicted values using graphs to evaluate forecasting accuracy visually.

Workflow of the Project

Project Workflow

Data Collection

↓

Data Preprocessing

↓

Data Exploration

↓

Data Visualization

↓

Stationarity Testing (ADF Test)

↓

ARIMA Model Development

↓

Temperature Forecasting

↓

Performance Evaluation

↓

Result Visualization

↓

Conclusion

Tools and Technologies

The following software tools and libraries were used for implementing the project.

Tool / Library	Purpose
Python	Programming Language
Pandas	Data Loading and Manipulation
NumPy	Numerical Computations
Matplotlib	Data Visualization
Statsmodels	Time Series Forecasting (ARIMA)
Scikit-learn	Performance Evaluation
Jupyter Notebook	Project Development Environment

Software Requirements

• Operating System : Windows 10 / Windows 11

• Programming Language : Python 3.x

• IDE : Jupyter Notebook

• Libraries : Pandas, NumPy, Matplotlib, Statsmodels, Scikit-learn

Hardware Requirements

• Processor : Intel Core i3 or Above

• RAM : Minimum 4 GB (8 GB Recommended)

• Storage : 500 MB Free Disk Space

• Internet Connection : Required for downloading the dataset and Python libraries

Data Preprocessing

Data preprocessing is one of the most important stages in any data science project. The quality of the input data directly affects the accuracy of the forecasting model.

The following preprocessing steps were performed:

• Imported the dataset into Python using Pandas.

• Checked the dataset structure and summary statistics.

• Verified that there were no missing values.

• Converted the Date column into datetime format.

• Set the Date column as the dataset index.

• Verified the chronological order of observations.

• Prepared the data for visualization and forecasting.

These preprocessing steps ensured that the dataset was clean, properly formatted, and ready for Time Series Analysis.

Time Series Analysis

Time Series Analysis is a statistical technique used to analyze data collected over regular intervals of time. Unlike traditional datasets where observations are treated as independent, time series data has a chronological order, and each observation may depend on previous observations.

The primary objective of time series analysis is to identify patterns, trends, seasonality, and temporal dependencies within historical data to forecast future values. Time series forecasting is widely used in weather prediction, stock market analysis, sales forecasting, traffic management, healthcare monitoring, and energy demand estimation.

A time series generally consists of the following components:

Trend

A trend represents the long-term movement in the data. It indicates whether the values are generally increasing, decreasing, or remaining constant over time.

Seasonality

Seasonality refers to repetitive patterns that occur at fixed intervals, such as daily, monthly, or yearly cycles.

Cyclic Pattern

Cyclic patterns are fluctuations that occur over longer periods and are influenced by external factors such as economic conditions or climate changes.

Random Noise

Random noise consists of unpredictable variations that cannot be explained by trends or seasonal patterns.

Understanding these components helps in selecting an appropriate forecasting model and improving prediction accuracy.

ARIMA Algorithm

ARIMA (AutoRegressive Integrated Moving Average) is one of the most widely used statistical models for time series forecasting. It is specifically designed to analyze historical observations and predict future values based on temporal relationships.

The ARIMA model is represented as:

ARIMA (p, d, q)

where:

• p = Number of AutoRegressive terms

• d = Degree of differencing required to make the series stationary

• q = Number of Moving Average terms

AutoRegressive (AR)

The AutoRegressive component predicts future values using previous observations. It assumes that past values have a direct influence on future values.

Integrated (I)

The Integrated component removes trends from the data by applying differencing. Differencing converts a non-stationary series into a stationary series, which is essential for accurate forecasting.

Moving Average (MA)

The Moving Average component models the relationship between the forecast errors and previous residual errors, helping to reduce forecasting inaccuracies.

The combination of these three components enables ARIMA to model a wide range of time series patterns effectively.

Stationarity Testing

Before applying the ARIMA model, it is necessary to verify whether the time series is stationary.

A stationary time series has constant statistical properties such as mean, variance, and covariance over time. Stationarity is important because ARIMA assumes that the underlying data does not exhibit changing statistical behavior.

The Augmented Dickey-Fuller (ADF) Test is used to determine stationarity.

Decision Rule:

• If the p-value is less than 0.05, the dataset is considered stationary.

• If the p-value is greater than 0.05, differencing is applied to make the dataset stationary.

In this project, the ADF test was performed before model development to ensure that the ARIMA algorithm could produce reliable forecasting results.

Model Implementation

The implementation of the forecasting model followed the steps below:

Step 1:
Import the required Python libraries, including Pandas, NumPy, Matplotlib, Statsmodels, and Scikit-learn.

Step 2:
Load the Daily Minimum Temperatures dataset into a Pandas DataFrame.

Step 3:
Convert the Date column into datetime format and use it as the dataset index.

Step 4:
Perform exploratory data analysis using statistical summaries and graphical visualizations.

Step 5:
Check for missing values and verify data quality.

Step 6:
Conduct the Augmented Dickey-Fuller (ADF) Test to determine stationarity.

Step 7:
Train the ARIMA model using historical temperature observations.

Step 8:
Generate future temperature forecasts.

Step 9:
Compare predicted temperatures with actual observations.

Step 10:
Evaluate the forecasting model using Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE).

Model Evaluation

The forecasting model is evaluated using two standard performance metrics.

Mean Absolute Error (MAE)

Mean Absolute Error calculates the average absolute difference between actual and predicted values.

A lower MAE indicates better prediction accuracy.

Root Mean Squared Error (RMSE)

Root Mean Squared Error measures the square root of the average squared prediction errors.

RMSE gives greater importance to larger prediction errors and is commonly used to evaluate forecasting models.

Lower RMSE values indicate that the forecasting model performs effectively.

Forecasting Process

The forecasting process begins by training the ARIMA model using historical temperature observations. After the model has learned the temporal patterns in the data, it generates predictions for future temperature values.

The predicted values are then compared with the actual observations to evaluate the model's performance. Graphical visualizations such as line charts are used to compare historical temperatures with forecasted temperatures.

The forecasting model can also be extended to predict future temperature values for longer periods, making it useful for weather forecasting and environmental analysis.

Overall, the forecasting process demonstrates how historical data can be transformed into valuable future predictions using advanced time series analysis techniques.

Results and Discussion

The Time Series Analysis model was successfully implemented using the Daily Minimum Temperatures dataset. The dataset was preprocessed, visualized, and analyzed before building the forecasting model. The ARIMA algorithm was then trained on the historical temperature data to predict future values.

Several visualizations were created during the analysis, including a temperature trend graph, histogram, box plot, rolling mean graph, rolling standard deviation graph, and actual versus predicted temperature graph. These visualizations helped in understanding the behavior of the dataset and evaluating the forecasting model.

The Augmented Dickey-Fuller (ADF) test was performed to verify the stationarity of the dataset before training the ARIMA model. After confirming stationarity, the model generated future temperature predictions based on historical observations.

The forecasting results showed that the predicted values closely followed the actual temperature values, demonstrating the effectiveness of the ARIMA model for time series forecasting. The model was evaluated using Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE), providing quantitative measures of forecasting accuracy.

Overall, the project successfully demonstrated the practical application of Time Series Analysis in weather forecasting and highlighted the importance of preprocessing, visualization, and model evaluation in achieving reliable predictions.

Challenges Faced

During the implementation of this project, several challenges were encountered.

Initially, understanding the concepts of time series forecasting and stationarity required additional study because time series data differs from traditional datasets. Selecting the appropriate ARIMA parameters also required experimentation to obtain better forecasting performance.

Managing date-based indexing and ensuring that the observations remained in chronological order were important preprocessing tasks. Visualizing long-term temperature trends and interpreting forecasting results also required careful analysis.

Despite these challenges, systematic preprocessing, exploratory data analysis, and model evaluation helped in building a reliable forecasting model.

Learning Outcomes

The completion of this project provided valuable practical experience in advanced data analysis techniques.

The following learning outcomes were achieved:

• Developed a clear understanding of Time Series Analysis.

• Learned the importance of stationarity in forecasting models.

• Gained practical experience in implementing the ARIMA algorithm.

• Improved skills in data preprocessing and feature preparation.

• Learned to visualize temporal data using Python libraries.

• Understood how to evaluate forecasting models using MAE and RMSE.

• Strengthened knowledge of Python libraries including Pandas, NumPy, Matplotlib, Statsmodels, and Scikit-learn.

• Gained experience in developing a complete end-to-end data science project.

Future Scope

The current project focuses on forecasting daily minimum temperatures using the ARIMA model. In the future, the project can be enhanced in several ways.

Advanced forecasting models such as SARIMA, Prophet, LSTM (Long Short-Term Memory), and GRU (Gated Recurrent Unit) can be implemented to improve prediction accuracy.

Additional weather parameters such as humidity, rainfall, wind speed, and atmospheric pressure can also be incorporated to develop multivariate forecasting models.

The project can further be extended into a real-time weather forecasting application by integrating live weather APIs and deploying the forecasting model as a web application.

Conclusion

This project successfully demonstrated the application of Time Series Analysis for forecasting daily minimum temperatures using historical weather data. The implementation covered every stage of the data science workflow, including data collection, preprocessing, visualization, stationarity testing, model development, forecasting, and evaluation.

The ARIMA model effectively captured the temporal patterns present in the dataset and generated reliable future temperature predictions. The forecasting results confirmed that time series analysis is a powerful analytical technique for solving real-world prediction problems.

This project not only enhanced practical knowledge of advanced data analysis techniques but also strengthened programming, data visualization, and machine learning skills. The experience gained from this project provides a strong foundation for applying forecasting techniques to other domains such as finance, healthcare, sales prediction, and energy demand forecasting.